

```
In [1]: # Importing all necessary libraries to create prediction model
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: # Loading csv file with DFW Airport snowfall data and setting DATE as index
dfw_data = '/Users/ali/Downloads/snowfall-dfw-airport.csv'
snow = pd.read_csv(dfw_data, index_col="DATE")

# Converting DATE index to datetime data type
snow.index = pd.to_datetime(snow.index)
print(snow.head())
```

| | STATION | NAME | SNOW | SNWD | TAVG |
|------------|-------------|------------------------------|------|------|------|
| DATE | | | | | |
| 2002-01-01 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 37.0 |
| 2002-01-02 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 28.0 |
| 2002-01-03 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 27.0 |
| 2002-01-04 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 38.0 |
| 2002-01-05 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 43.0 |

```
In [3]: # Verifying the details of the dataset (data types, etc.)
print(snow.info())
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 7306 entries, 2002-01-01 to 2022-01-01
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   STATION     7306 non-null   object
 1   NAME        7306 non-null   object
 2   SNOW        7306 non-null   float64
 3   SNWD        5997 non-null   float64
 4   TAVG        4502 non-null   float64
dtypes: float64(3), object(2)
memory usage: 342.5+ KB
None
```

```
In [4]: # Looking for null values by percent in each column
snow.apply(pd.isnull).sum() / snow.shape[0]
```

```
Out[4]: STATION      0.000000
NAME              0.000000
SNOW              0.000000
SNWD              0.179168
TAVG              0.383794
dtype: float64
```

```
In [5]: # Checking for null values in Snow Depth (SNWD) and Temp Average (TAVG) col
snow[pd.isnull(snow["SNWD"])]
```

Out[5]:

| | STATION | NAME | SNOW | SNWD | TAVG |
|------------|-------------|------------------------------|------|------|------|
| DATE | | | | | |
| 2002-01-01 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 37.0 |
| 2002-01-02 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 28.0 |
| 2002-01-03 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 27.0 |
| 2002-01-04 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 38.0 |
| 2002-01-05 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | 43.0 |
| ... | ... | ... | ... | ... | ... |
| 2005-11-26 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | NaN |
| 2005-11-27 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | NaN |
| 2005-11-28 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | NaN |
| 2005-11-29 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | NaN |
| 2005-11-30 | USW00003927 | DAL FTW WSCMO AIRPORT, TX US | 0.0 | NaN | NaN |

1309 rows × 5 columns

```
In [6]: # Replacing null values in Snow Depth column with 0
snow["SNWD"] = snow["SNWD"].fillna(0.0)
```

```
In [7]: # Making sure we have no more null values in Snow Depth column
snow[pd.isnull(snow["SNWD"])]
```

Out[7]:

| STATION | NAME | SNOW | SNWD | TAVG |
|---------|------|------|------|------|
| DATE | | | | |

```
In [8]: # Using Pandas "foward fill" feature take fill Temp Average with last non-n
snow = snow.fillna(method="ffill")
```

```
In [9]: # Making sure we have no more null values in Temp Average column
snow[pd.isnull(snow["TAVG"])]
```

Out[9]:

| STATION | NAME | SNOW | SNWD | TAVG |
|---------|------|------|------|------|
| DATE | | | | |


```
In [10]: # Deleting non-numeric columns to keep the data clean
del snow["STATION"]
del snow["NAME"]
```

```
In [11]: # Checking to see the updated number of rows in dataset
print(snow.info())
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 7306 entries, 2002-01-01 to 2022-01-01
Data columns (total 3 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   SNOW    7306 non-null     float64
 1   SNWD    7306 non-null     float64
 2   TAVG    7306 non-null     float64
dtypes: float64(3)
memory usage: 228.3 KB
None
```


```
In [12]: # Plotting the snow values with snow depth
snow[["SNOW", "SNWD"]].plot()
```

```
Out[12]: <AxesSubplot:xlabel='DATE'>
```



```
In [13]: # Plotting the temp average values to see correlation with snowfall  
snow["TAVG"].plot()
```

```
Out[13]: <AxesSubplot:xlabel='DATE'>
```



```
In [14]: # Pulling the amount of snowfall by year
snow.groupby(snow.index.year).sum() [ "SNOW" ]
```

```
Out[14]: DATE
2002      3.8
2003      3.2
2004      2.9
2005      0.0
2006      0.0
2007      0.3
2008      2.1
2009      3.4
2010     13.9
2011      4.3
2012      0.8
2013      2.4
2014      0.8
2015      5.8
2016      0.0
2017      0.1
2018      0.0
2019      0.0
2020      0.2
2021      5.0
2022      0.0
Name: SNOW, dtype: float64
```

```
In [15]: # Adding a prediction column for snowfall, 1 day in the future
# Using the 'shift' function to move all data back by 1 row
snow[ "PREDICTION" ] = snow.shift(-1)[ "SNOW" ]
```

```
In [16]: # Copying all rows except last one with null value in "Prediction" column
snow = snow.iloc[:-1,:].copy()
```


```
In [17]: snow.describe()
```

```
Out[17]:
```

| | SNOW | SNWD | TAVG | PREDICTION |
|-------|-------------|-------------|-------------|-------------|
| count | 7305.000000 | 7305.000000 | 7305.000000 | 7305.000000 |
| mean | 0.006708 | 0.009938 | 73.699658 | 0.006708 |
| std | 0.167610 | 0.196512 | 14.807791 | 0.167610 |
| min | 0.000000 | 0.000000 | 8.000000 | 0.000000 |
| 25% | 0.000000 | 0.000000 | 64.000000 | 0.000000 |
| 50% | 0.000000 | 0.000000 | 83.000000 | 0.000000 |
| 75% | 0.000000 | 0.000000 | 84.000000 | 0.000000 |
| max | 11.200000 | 11.000000 | 96.000000 | 11.200000 |


```
In [18]: sns.pairplot(snow)
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x7fd4748076d0>
```



```
In [19]: snow.plot(x='SNOW', y='TAVG', kind = 'scatter')
```

```
Out[19]: <AxesSubplot:xlabel='SNOW', ylabel='TAVG'>
```



```
In [20]: # Importing the Ridge regression library
from sklearn.linear_model import Ridge

# Setting up the L2 penalty for inputs that hinder the prediction
reg_model = Ridge(alpha=.1)
```

```
In [21]: # Setting up the predictors for next day's weather
predictors = ["SNOW", "TAVG"]
```

```
In [22]: # Setting up our train and test sets
train = snow.loc[:"2020-12-31"]
test = snow.loc["2021-01-01":]
```

```
In [23]: reg_model.fit(train[predictors], train["PREDICTION"])
```

```
Out[23]: Ridge(alpha=0.1)
```

```
In [24]: # Calculate predictions based on regression model with test set
predictions = reg_model.predict(test[predictors])
```

```
In [25]: # Importing the mean_absolute_error library to evaluate predictions
from sklearn.metrics import mean_absolute_error
mean_absolute_error(test["PREDICTION"], predictions)
```

Out[25]: 0.02187426057605117


```
In [26]: # Looking at the results based on actual snowfall vs predicted snowfall
results = pd.concat([test["PREDICTION"], pd.Series(predictions, index=test.
results.columns = ["ACTUAL", "PREDICTION"]
print(results)
```

| | ACTUAL | PREDICTION |
|------------|--------|------------|
| DATE | | |
| 2021-01-01 | 0.0 | 0.013170 |
| 2021-01-02 | 0.0 | 0.012168 |
| 2021-01-03 | 0.0 | 0.011166 |
| 2021-01-04 | 0.0 | 0.010565 |
| 2021-01-05 | 0.0 | 0.010565 |
| ... | ... | ... |
| 2021-12-27 | 0.0 | 0.006558 |
| 2021-12-28 | 0.0 | 0.006157 |
| 2021-12-29 | 0.0 | 0.007560 |
| 2021-12-30 | 0.0 | 0.008161 |
| 2021-12-31 | 0.0 | 0.008161 |

[365 rows x 2 columns]

```
In [27]: results.plot()
```

Out[27]: <AxesSubplot:xlabel='DATE'>




```
In [28]: # Evaluating the difference between the actual snowfall vs predicted snowfa
results["DIFFERENCE"] = (results["ACTUAL"] - results["PREDICTION"]).abs()
results.sort_values("ACTUAL", ascending=False).head(25)
```

Out[28]:

| | ACTUAL | PREDICTION | DIFFERENCE |
|------------|--------|------------|------------|
| DATE | | | |
| 2021-02-13 | 4.0 | 0.015774 | 3.984226 |
| 2021-02-16 | 1.0 | 0.018980 | 0.981020 |
| 2021-08-29 | 0.0 | 0.004154 | 0.004154 |
| 2021-09-07 | 0.0 | 0.004154 | 0.004154 |
| 2021-09-06 | 0.0 | 0.003753 | 0.003753 |
| 2021-09-05 | 0.0 | 0.003352 | 0.003352 |
| 2021-09-04 | 0.0 | 0.003152 | 0.003152 |
| 2021-09-03 | 0.0 | 0.003152 | 0.003152 |
| 2021-09-02 | 0.0 | 0.002951 | 0.002951 |
| 2021-09-01 | 0.0 | 0.002751 | 0.002751 |
| 2021-08-31 | 0.0 | 0.002951 | 0.002951 |
| 2021-08-30 | 0.0 | 0.003552 | 0.003552 |
| 2021-08-28 | 0.0 | 0.003552 | 0.003552 |
| 2021-08-18 | 0.0 | 0.005556 | 0.005556 |
| 2021-08-27 | 0.0 | 0.003552 | 0.003552 |
| 2021-08-26 | 0.0 | 0.002751 | 0.002751 |
| 2021-08-25 | 0.0 | 0.003152 | 0.003152 |
| 2021-08-24 | 0.0 | 0.003152 | 0.003152 |
| 2021-08-23 | 0.0 | 0.003152 | 0.003152 |
| 2021-08-22 | 0.0 | 0.003152 | 0.003152 |
| 2021-08-21 | 0.0 | 0.003152 | 0.003152 |
| 2021-08-20 | 0.0 | 0.003152 | 0.003152 |
| 2021-09-08 | 0.0 | 0.004154 | 0.004154 |
| 2021-09-09 | 0.0 | 0.004554 | 0.004554 |
| 2021-09-10 | 0.0 | 0.004554 | 0.004554 |

